

Why Shadow Volumes?


- Dynamic shadows improve your game
 - Dramatic effects
 - Better sense of 3D
- Most game shadows today are very limited
 - Planar projected shadows [Blinn '88]
 - Limited to floor planes, perhaps walls



Not good enough
for today's games!




Stenciled Shadow Volumes in Practice



Notice the proper
self-shadowing!



2D Cutaway of a Shadow Volume



Shadow Volume Advantages

- Omni-directional approach
 - Not just spotlight frustums as with shadow maps
- Automatic self-shadowing
 - Everything can shadow everything, including self
 - Without *shadow acne* artifacts as with shadow maps
- Window-space shadow determination
 - Shadows accurate to a pixel
 - Or sub-pixel if multisampling is available
- Required stencil buffer broadly supported today
 - OpenGL support since version 1.0 (1991)
 - Direct3D support since DX6 (1998)




Shadow Volume Disadvantages

- Ideal light sources only
 - Limited to local point and directional lights
 - No area light sources for soft shadows
- Requires polygonal models with connectivity
 - Models must be closed (2-manifold)
 - Models must be free of non-planar polygons
- Silhouette computations are required
 - Can burden CPU
 - Particularly for dynamic scenes
- Inherently multi-pass algorithm
- Consumes lots of GPU fill rate



Visualizing Shadow Volumes in 3D

- Occluders and light source cast out a shadow volume
- Objects within the volume should be shadowed




Scene with shadows from an NVIDIA logo casting a shadow volume

Visualization of the shadow volume




Visualizing the Stencil Buffer Counts

Shadowed scene



Stencil buffer contents



red = stencil value of 1
green = stencil value of 0

Stencil counts beyond 1 are possible for multiple or complex occluders.

GLUT *shadowvol* example credit: Tom McReynolds




Counting Shadow Volume Enter/Leaves With a Stencil Buffer (Zpass approach)


- Render scene to initialize depth buffer
 - Depth values indicate the closest visible fragments
- Use a stencil enter/leave counting approach
 - Draw shadow volume twice using face culling
 - 1st pass: render front faces and increment when depth test passes
 - 2nd pass: render back faces and decrement when depth test passes
 - Don't update depth or color
- Afterward, pixel's stencil is non-zero if pixel in shadow, and zero if illuminated



Why Eye-to-Object Stencil Enter/Leave Counting Approach Works




Illuminated, Behind Shadow Volumes



Shadow Volume Count = $+1+1+1-1-1-1 = 0$




Shadowed, Nested in Shadow Volumes




Shadow Volume Count = $+1+1+1-1 = 2$



Illuminated, In Front of Shadow Volumes



Problems Created by Near Plane Clipping (Zpass approach)



Alternative Approach: Zfail

- ❏ Render scene to initialize depth buffer
 - ❏ Depth values indicate the closest visible fragments
- ❏ Use a stencil enter/leave counting approach
 - ❏ Draw shadow volume twice using face culling
 - ❏ 1st pass: render back faces and increment when depth test fails
 - ❏ 2nd pass: render front faces and decrement when depth test fails
 - ❏ Don't update depth or color
- ❏ Afterward, pixel's stencil is non-zero if pixel in shadow, and zero if illuminated

Zfail versus Zpass Comparison (1)

- ❏ When stencil increment/decrements occur
 - ❏ Zpass: on depth test pass
 - ❏ Zfail: on depth test fail
- ❏ Increment on
 - ❏ Zpass: front faces
 - ❏ Zfail: back faces
- ❏ Decrement on
 - ❏ Zpass: front faces
 - ❏ Zfail: back faces

Zfail versus Zpass Comparison (2)

- ❏ Both cases order passes based stencil operation
 - ❏ First, render increment pass
 - ❏ Second, render decrement pass
 - ❏ Why?
 - ❏ Because standard stencil operations saturate
 - ❏ Wrapping stencil operations can avoid this
- ❏ Which clip plane creates a problem
 - ❏ Zpass: near clip plane
 - ❏ Zfail: far clip plane
- ❏ Either way is foiled by view frustum clipping
 - ❏ Which clip plane (front or back) changes

Insight!

- ❏ If we could avoid *either* near plane or far plane view frustum clipping, shadow volume rendering could be robust
- ❏ Avoiding near plane clipping
 - ❏ Not really possible
 - ❏ Objects can always be behind you
 - ❏ Moreover, depth precision in a perspective view goes to hell when the near plane is too near the eye
- ❏ Avoiding far plane clipping
 - ❏ Perspective make it possible to render at infinity
 - ❏ Depth precision is terrible at infinity, but we just care about avoiding clipping

Avoiding Far Plane Clipping

- Usual practice for perspective GL projection matrix
 - Use *glFrustum* (or *gluPerspective*)
 - Requires two values for near & far clip planes
 - Near plane's distance from the eye
 - Far plane's distance from the eye
 - Assumes a *finite* far plane distance
- Alternative projection matrix
 - Still requires near plane's distance from the eye
 - But assume far plane is at *infinity*
- What is the limit of the projection matrix when the far plane distance goes to infinity?



Standard *glFrustum* Projection Matrix

$$P = \begin{bmatrix} \frac{2 \times \text{Near}}{\text{Right} - \text{Left}} & 0 & \frac{\text{Right} + \text{Left}}{\text{Right} - \text{Left}} & 0 \\ 0 & \frac{2 \times \text{Near}}{\text{Top} - \text{Bottom}} & \frac{\text{Top} + \text{Bottom}}{\text{Top} - \text{Bottom}} & 0 \\ 0 & 0 & \frac{\text{Far} + \text{Near}}{\text{Far} - \text{Near}} & -\frac{2 \times \text{Far} \times \text{Near}}{\text{Far} - \text{Near}} \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

- Only third row depends on *Far* and *Near*



Limit of *glFrustum* Projection Matrix as Far Plane is Moved to Infinity

$$\lim_{\text{Far} \rightarrow \infty} P = P_{\text{inf}} = \begin{bmatrix} \frac{2 \times \text{Near}}{\text{Right} - \text{Left}} & 0 & \frac{\text{Right} + \text{Left}}{\text{Right} - \text{Left}} & 0 \\ 0 & \frac{2 \times \text{Near}}{\text{Top} - \text{Bottom}} & \frac{\text{Top} + \text{Bottom}}{\text{Top} - \text{Bottom}} & 0 \\ 0 & 0 & -1 & -2 \times \text{Near} \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

- First, second, and fourth rows are the same as in *P*
- But third row *no longer* depends on *Far*
 - Effectively, *Far* equals ∞



Verifying P_{inf} Will Not Clip Infinitely Far Away Vertices (1)

- What is the most distant possible vertex in front of the eye?
 - Ok to use homogeneous coordinates
 - OpenGL convention looks down the negative Z axis
 - So most distant vertex is (0,0,-D,0) where D>0
- Transform (0,0,-D,0) to window space
 - Is such a vertex clipped by P_{inf} ?
 - No, it is not clipped, as explained on the next slide



Verifying P_{inf} Will Not Clip Infinitely Far Away Vertices (2)

- Transform eye-space (0,0,-D,0) to clip-space

$$\begin{bmatrix} x_c \\ y_c \\ -D \\ -D \end{bmatrix} = \begin{bmatrix} x_e \\ y_e \\ z_e \\ w_e \end{bmatrix} = \begin{bmatrix} \frac{2 \times \text{Near}}{\text{Right} - \text{Left}} & 0 & \frac{\text{Right} + \text{Left}}{\text{Right} - \text{Left}} & 0 \\ 0 & \frac{2 \times \text{Near}}{\text{Top} - \text{Bottom}} & \frac{\text{Top} + \text{Bottom}}{\text{Top} - \text{Bottom}} & 0 \\ 0 & 0 & -1 & -2 \times \text{Near} \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ -D \\ 0 \end{bmatrix}$$

- Then, assuming *glDepthRange*(0,1), transform clip-space position to window-space position

$$z_w = 0.5 \times \frac{z_c}{w_c} + 0.5 = 0.5 \times \frac{-D}{-D} + 0.5 = 1$$

- So ∞ in front of eye transforms to the maximum window-space Z value, but is still within the valid depth range (i.e., not clipped)



Is P_{inf} Bad for Depth Buffer Precision?

- Naïve question
 - Wouldn't moving the far clip plane to infinity waste depth buffer precision? Seems plausible, but
- Answer: Not really
 - Minimal depth buffer precision is wasted in practice
 - This is due to projective nature of perspective
- Say *Near* is 1.0 and *Far* is 100.0 (typical situation)
 - P* would transform eye-space infinity to only 1.01 in window space
 - Only a 1% compression of the depth range is required to render infinity without clipping
 - Moving near closer would hurt precision



P_{inf} Depth Precision Scale Factor

- Using P_{inf} with Near instead of P with Near and Far compresses (scales) the depth precision by

$$\frac{(Far - Near)}{Far}$$

- The compression of depth precision is uniform, but the depth precision itself is already non-uniform on eye-space interval [Near, Far] due to perspective
 - So the discrete loss of precision is more towards the far clip plane
- Normally, $Far \gg Near$ so the scale factor is usually less than but still nearly 1.0
 - So the compression effect is minor



Robust Stenciled Shadow Volumes Without Near (or Far) Plane Capping

- Use Zfail Stenciling Approach
 - Must render geometry to close shadow volume extrusion on the model and at infinity (explained later)
- Use the P_{inf} Projection Matrix
 - No worries about far plane clipping
 - Losses some depth buffer precision (but not much)
- Draw the infinite vertices of the shadow volume using homogeneous coordinates (w=0)



Rendering Closed, but Infinite, Shadow Volumes

- To be robust, the shadow volume geometry must be closed, even at infinity
- Three sets of polygons close the shadow volume
 - Possible silhouette edges extruded to infinity away from the light
 - All of the occluder's back-facing (w.r.t. the light) triangles projected away from the light to infinity
 - All of the occluder's front-facing (w.r.t. the light) triangles
- We assume the object vertices and light position are homogeneous coordinates, i.e. (x,y,z,w)
 - Where $w \geq 0$



1st Set of Shadow Volume Polygons

- Assuming
 - A and B are vertices of an occluder model's possible silhouette edge
 - And L is the light position
- For all A and B on silhouette edges of the occluder model, render the quad

$$\begin{aligned} &\langle B_x, B_y, B_z, B_w \rangle \\ &\langle A_x, A_y, A_z, A_w \rangle \\ &\langle A_x L_w - L_x A_w, A_y L_w - L_y A_w, A_z L_w - L_z A_w, 0 \rangle \\ &\langle B_x L_w - L_x B_w, B_y L_w - L_y B_w, B_z L_w - L_z B_w, 0 \rangle \end{aligned}$$

$\swarrow \searrow$
 Homogenous
vector differences
- What is a possible silhouette edge?
 - One polygon sharing an edge faces toward L
 - Other faces away from L



Examples of Possible Silhouette Edges for Quake2 Models



An object viewed from the same basic direction that the light is shining on the object has an identifiable light-view silhouette

An object's light-view silhouette appears quite jumbled when viewed from a point-of-view that does not correspond well with the light's point-of-view



2nd and 3rd Set of Shadow Volume Polygons

- 2nd set of polygons
 - Assuming A, B, and C are each vertices of occluder model's back-facing triangles w.r.t. light position L

$$\begin{aligned} &\langle A_x L_w - L_x A_w, A_y L_w - L_y A_w, A_z L_w - L_z A_w, 0 \rangle \\ &\langle B_x L_w - L_x B_w, B_y L_w - L_y B_w, B_z L_w - L_z B_w, 0 \rangle \\ &\langle C_x L_w - L_x C_w, C_y L_w - L_y C_w, C_z L_w - L_z C_w, 0 \rangle \end{aligned}$$

$\swarrow \searrow$
 Homogenous
vector differences
 - These vertices are effectively directions (w=0)
- 3rd set of polygons
 - Assuming A, B, and C are each vertices of occluder model's front-facing triangles w.r.t. light position L

$$\begin{aligned} &\langle A_x, A_y, A_z, A_w \rangle \\ &\langle B_x, B_y, B_z, B_w \rangle \\ &\langle C_x, C_y, C_z, C_w \rangle \end{aligned}$$

$\swarrow \searrow$
 Homogenous
vector differences



Complete Robust Stenciled Shadow Volume Rendering Technique

- See our paper “Practical and Robust Stenciled Shadow Volumes for Hardware-Accelerated Rendering”
 - In the accompanying course notes
 - And on-line at developer.nvidia.com
- Paper has pseudo-code for rendering procedure
 - OpenGL state settings & rendering commands
 - Supports multiple per-vertex lights
 - Assumes application computes object-space determination of occluder model's polygons orientation w.r.t. each light



Requirements for Our Stenciled Shadow Volume Technique (1)

- Models must be composed of triangles only (avoiding non-planar polygons)
- Models must be closed (2-manifold) and have a consistent winding order
 - Bergeron ['86] approach could be used to handle “open” models if necessary
- Homogeneous object coordinates are permitted, assuming $w \geq 0$
 - If not, $(x, y, z, -1) = (-x, -y, -z, 1)$
- Ideal light sources only
 - Directional or positional, assuming $w \geq 0$



Requirements for Our Stenciled Shadow Volume Technique (2)

- Connectivity information for occluding models must be available
 - So silhouette edges w.r.t. light positions can be determined at shadow volume construction time
- Projection matrix must be perspective
 - Not orthographic
 - NV_depth_clamp extension provides orthographic support (more later)
- Render must guarantee “watertight” rasterization
 - No double hitting pixels at shared polygon edges
 - No missed pixels at shared polygon edges



Requirements for Our Stenciled Shadow Volume Technique (3)

- Enough stencil bits
 - N stencil bits where 2^N is greater than the maximum shadow depth count ever encountered
 - Scene dependent
 - 8-bits is usually quite adequate & what all recent stencil hardware provides
 - Wrapping stencil increment/decrement operations (i.e. OpenGL's EXT_stencil_wrap) permit deeper shadow counts, modulo aliasing with zero
 - Realize that shadow depths > 255 imply too much fill rate for interactive applications



Requirements for Our Stenciled Shadow Volume Technique (4)

- Rendering features provided by OpenGL 1.0 or DirectX 6 (or subsequent versions)
 - Transformation & clipping of homogenous positions
 - Front- and back-face culling
 - Masking color and depth buffer writes
 - Depth buffering (i.e. conventional Z-buffering)
 - Stencil-testing support

In practice, these are quite reasonable requirements for nearly any polygonal-based 3D game or application



Our Approach in Practice (1)



Scene with shadows.
Yellow light is embedded in the green three-holed object.
 P_{inf} is used for all the following scenes.

Same scene visualizing the shadow volumes.



Our Approach in Practice (2)

Details worth noting . . .




Fine details: Shadows of the A, N, and T letters on the knight's armor and shield.




Hard case: The shadow volume from the front-facing hole would definitely intersect the near clip plane.



Our Approach in Practice (3)




Alternate view of same scene with shadows. Yellow lines indicate previous view's view frustum boundary. Recall shadows are view-independent.




Shadow volumes from the alternate view.



Our Approach in Practice (4)



Clip-space view. Original view's scene seen from clip space. The back plane is "at infinity" with very little effective depth precision near infinity.



Clip-space view of shadow volumes. Back-facing triangles w.r.t. light are seen projected onto far plane at infinity.



Another Example (1)




Original eye's view. Again, yellow light is embedded in the green three-holed object. P_{inf} is used for all the following scenes.



Eye-space view of previous eye's view. Clipped to the previous eye's P_{inf} view frustum. Shows knight's projection to infinity.



Another Example (2)



Clip-space view of previous eye's view. Shows shadow volume closed at infinity and other shadow volume's intersection with the near clip plane.

Original eye's far clip plane

Original eye's near clip plane



Stencil Shadow Volumes with Multiple Lights



Three colored lights. Diffuse/specular bump mapped animated characters with shadows. 34 fps on GeForce4 Ti 4600; 80+ fps for one light.



Stencil Shadow Volumes for Simulating Soft Shadows



Cluster of 12 dim lights approximating an area light source. Generates a soft shadow effect; careful about banding. 8 fps on GeForce4 Ti 4600.

The cluster of point lights.



Stencil Shadow Volume Optimizations (1)

- ❏ Use `GL_QUAD_STRIP` rather than `GL_QUADS` to render extruded shadow volume quads
 - ❏ Requires determining possible silhouette loop connectivity
- ❏ Mix *Zfail* and *Zpass* techniques
 - ❏ Pick a single formulation for each shadow volume
 - ❏ *Zpass* is more efficient since shadow volume does not need to be closed
 - ❏ Mixing has no effect on net shadow depth count
 - ❏ *Zfail* can be used in the hard cases



Stencil Shadow Volume Optimizations (2)

- ❏ Pre-compute or re-use cache shadow volume geometry when geometric relationship between a light and occluder does not change between frames
 - ❏ Example: Headlights on a car have a static shadow volume w.r.t. the car itself as an occluder
- ❏ Advanced shadow volume culling approaches
 - ❏ Uses portals, Binary Space Partitioning trees, occlusion detection, and view frustum culling techniques to limit shadow volumes
 - ❏ Careful to make sure such optimizations are truly correct



Stencil Shadow Volume Optimizations (3)

- ❏ Take advantage of ad-hoc knowledge of scenes whenever possible
 - ❏ Example: A totally closed room means you do not have to cast shadow volumes for the wall, floor, ceiling
- ❏ Limit shadows to important entities
 - ❏ Example: Generate shadow volumes for monsters and characters, but not static objects
 - ❏ Characters can still cast shadows on static objects
- ❏ Mix shadow techniques where possible
 - ❏ Use planar projected shadows or light-maps where appropriate




Stencil Shadow Volume Optimizations (4)

- ❏ Shadow volume's extrusion for directional lights can be rendered with a `GL_TRIANGLE_FAN`
 - ❏ Directional light's shadow volume always projects to a single point at infinity



Scene with directional light.



Clip-space view of shadow volume



Hardware Enhancements: Wrapping Stencil Operations

- ❏ Conventional OpenGL 1.0 stencil operations
 - ❏ `GL_INCR` increments and clamps to $2^N - 1$
 - ❏ `GL_DECR` decrements and clamps to zero
- ❏ DirectX 6 introduced "wrapping" stencil operations
- ❏ Exposed by OpenGL's `EXT_stencil_wrap` extension
 - ❏ `GL_INCR_WRAP_EXT` increments modulo 2^N
 - ❏ `GL_DECR_WRAP_EXT` decrements modulo 2^N
- ❏ Avoids saturation throwing off the shadow volume depth count
 - ❏ Still possible, though very rare, that 2^N , 2×2^N , 3×2^N , etc. can alias to zero



Hardware Enhancements: Depth Clamp (1)

- What is depth clamping?
 - Boolean hardware enable/disable
 - When enabled, disables the near & far clip planes
 - Interpolate the window-space depth value
 - Clamps the interpolated depth value to the depth range, i.e. $[\min(n, f), \max(n, f)]$
 - Assuming `glDepthRange(n, f);`
 - Geometry "behind" the far clip plane is still rendered
 - Depth value clamped to farthest Z
 - Similar for near clip plane, as long as $w > 0$, except clamped to closest Z



Hardware Enhancements: Depth Clamp (2)

- Advantage for stenciled shadow volumes
 - With depth clamp, P (rather than P_{in}) can be used with our robust stenciled shadow volume technique
 - Marginal loss of depth precision re-gained
 - Orthographic projections can work with our technique with depth clamping
- NV_depth_clamp OpenGL extension
 - Easy to use
 - `glEnable(GL_DEPTH_CLAMP_NV);`
 - `glDisable(GL_DEPTH_CLAMP_NV);`
 - GeForce3 & GeForce4 Ti support (soon)



Hardware Enhancements: Two-sided Stencil Testing (1)

- Current stenciled shadow volumes required rendering shadow volume geometry twice
 - First, rasterizing front-facing geometry
 - Second, rasterizing back-facing geometry
- Two-sided stencil testing requires only one pass
 - Two sets of stencil state: front- and back-facing
 - Boolean enable for two-sided stencil testing
 - When enabled, back-facing stencil state is used for stencil testing back-facing polygons
 - Otherwise, front-facing stencil state is used
 - Rasterizes just as many fragments, but more efficient for CPU & GPU



Hardware Enhancements: Two-sided Stencil Testing (2)

- NV_stencil_two_side OpenGL extension
 - Enable applies if `GL_STENCIL_TEST` also enabled
 - `glEnable(GL_STENCIL_TEST_TWO_SIDE_NV);`
 - `glDisable(GL_STENCIL_TEST_TWO_SIDE_NV);`
 - Control of front- and back-facing stencil state update
 - `glActiveStencilFaceNV(GL_FRONT);`
 - `glActiveStencilFaceNV(GL_BACK);`
 - Existing stencil routines (`glStencilOp`, `glStencilMask`, `glStencilFunc`) update the active stencil face state
 - `glClear` and non-polygon primitives always use the front-facing stencil state
 - Expect on future GPUs



Usage of NV_stencil_two_side & EXT_stencil_wrap

OLD SCHOOL

```
glDepthMask(0);
glColorMask(0,0,0,0);
glEnable(GL_CULL_FACE);
glEnable(GL_STENCIL_TEST);
glStencilMask(-0);
glStencilFunc(GL_ALWAYS, 0, -0);
// Increment for back faces
glCullFace(GL_BACK);
glStencilOp(GL_KEEP, // stencil test fail
           GL_INCR, // depth test fail
           GL_INCR); // depth test pass
renderShadowVolumePolygons();
// Decrement for front faces
glCullFace(GL_FRONT);
glStencilOp(GL_KEEP, // stencil test fail
           GL_DECR, // depth test fail
           GL_DECR); // depth test pass
renderShadowVolumePolygons();
```

New approach calls `renderShadowVolumePolygons()` just once.



NEW SCHOOL

```
glDepthMask(0);
glColorMask(0,0,0,0);
glDisable(GL_CULL_FACE);
glDisable(GL_STENCIL_TEST);
glEnable(GL_STENCIL_TEST_TWO_SIDE_NV);
glActiveStencilFaceNV(GL_BACK);
glStencilOp(GL_KEEP, // stencil test fail
           GL_INCR_WRAP_EXT, // depth test fail
           GL_KEEP); // depth test pass
glStencilMask(-0);
glStencilFunc(GL_ALWAYS, 0, -0);
glActiveStencilFaceNV(GL_FRONT);
glStencilOp(GL_KEEP, // stencil test fail
           GL_DECR_WRAP_EXT, // depth test fail
           GL_KEEP); // depth test pass
glStencilMask(-0);
glStencilFunc(GL_ALWAYS, 0, -0);
renderShadowVolumePolygons();
```

Shadow Volume History (1)

- Invented by Frank Crow ['77]
 - Software rendering scan-line approach
- Brotman and Badler ['84]
 - Software-based depth-buffered approach
 - Used lots of point lights to simulate soft shadows
- Pixel-Planes [Fuchs, et.al. '85] hardware
 - First hardware approach
 - Point within a volume, rather than ray intersection
- Bergeron ['96] generalizations
 - Explains how to handle open models
 - And non-planar polygons



Shadow Volume History (2)

- ❏ Fournier & Fussell ['88] theory
 - ❏ Provides theory for shadow volume counting approach within a frame buffer
- ❏ Akeley & Foran invent the stencil buffer
 - ❏ IRIS GL functionality, later made part of OpenGL 1.0
 - ❏ Patent filed in '92
- ❏ Heidmann [*IRIS Universe article*, '91]
 - ❏ IRIS GL stencil buffer-based approach
- ❏ Deifenbach's thesis ['96]
 - ❏ Used stenciled volumes in multi-pass framework



Shadow Volume History (3)

- ❏ Dietrich slides [March '99] at GDC
 - ❏ Proposes *zfail* based stenciled shadow volumes
- ❏ Kilgard whitepaper [March '99] at GDC
 - ❏ *Invert* approach for planar cut-outs
- ❏ Bilodeau slides [May '99] at Creative seminar
 - ❏ Proposes way around near plane clipping problems
 - ❏ Reverses depth test function to reverse stencil volume ray intersection sense
- ❏ Carmack [unpublished, early 2000]
 - ❏ First detailed discussion of the equivalence of *zpass* and *zfail* stenciled shadow volume methods



Shadow Volume History (4)

- ❏ Kilgard [2001] at GDC and CEDEC Japan
 - ❏ Proposes *zpass* capping scheme
 - ❏ Project back-facing (w.r.t. light) geometry to the near clip plane for capping
 - ❏ Establishes *near plane ledge* for crack-free near plane capping
 - ❏ Applies homogeneous coordinates ($w=0$) for rendering infinite shadow volume geometry
 - ❏ Requires much CPU effort for capping
 - ❏ Not totally robust because CPU and GPU computations will not match exactly, resulting in cracks



Our Contribution

- ❏ We Integrate Multiple Ideas into a Robust Solution
 - ❏ Dietrich and Carmack's *zfail* approach
 - ❏ Kilgard's homogeneous coordinates ($w=0$) for rendering infinite shadow volume geometry
 - ❏ Somewhat-obscure [Blinn '93] *infinite far plane* projection matrix formulation
 - ❏ DirectX 6's wrapping stencil increment & decrement
 - ❏ OpenGL's EXT_stencil_wrap extension
- ❏ Hardware Enhancements
 - ❏ Depth clamping: *better depth precision*
 - ❏ Two-sided stencil testing: *performance*



Conclusions

- ❏ Presented robust stenciled shadow volume technique
 - ❏ Reasonable & well-defined requirements
 - ❏ Almost impossible to buy PC in 2002 without stencil buffer support
 - ❏ Still, shadow volumes require high fill rates
- ❏ Technical paper has more detail
- ❏ Carmack's new Doom engine is already using this technique today
 - ❏ Why aren't you?

